

Microservices with Quarkus

Day 3: Building Resilient & Asynchronous Microservices

Morning Session 1 (09:00 - 10:30)

Topic: Building Resilient Microservices

Activity: Hands-on Lab 6

Lab 6 Objectives

- Apply resilience patterns to the `train-line-service` .
- Handle latency and transient failures from the `station-service` .
- Use `@Timeout` and `@Retry` to meet a 90% success rate objective.
- Implement a `@CircuitBreaker` to prevent cascading failures.
- Use a `@Fallback` to provide a clean response when failures occur.

What is Resilience?

In a distributed system, failures are inevitable.

- Network issues
- Service unavailability
- High latency

Resilience is the ability of a system to gracefully handle these failures and remain functional.

MicroProfile Fault Tolerance

Provides a standard set of annotations to build resilient microservices.

- **Goal:** Isolate and manage failures, preventing them from cascading.
- **Key Annotations:**
 - `@Timeout` : Sets a duration limit for a method execution.
 - `@Retry` : Automatically retries a failed execution.
 - `@CircuitBreaker` : Stops repeated calls to a failing service.
 - `@Fallback` : Provides an alternative result for a failed execution.
 - `@Bulkhead` : Limits concurrent requests to a method.
 - `@RateLimit` : Limit the number of requests over time.
- **Guide:** [Fault Tolerance with SmallRye](#)

SmallRye Fault Tolerance

- **Implementation:** Quarkus uses SmallRye to implement the MicroProfile Fault Tolerance specification.
- **Beyond the Spec:** SmallRye provides additional features and fine-grained configuration.
- **Asynchronous Support:** All fault tolerance annotations work with both synchronous and asynchronous (`CompletionStage` , `Uni`) methods.
- **Configuration:** While annotations are easy, all parameters can be configured and overridden in `application.properties` for greater flexibility.

@Timeout

- **Problem:** A synchronous call is taking too long, tying up resources.
- **Solution:** Abort the operation if it exceeds a specified duration.

```
@GET
@Path("/{id}")
@Timeout(250) // Fails if execution takes longer than 250ms
public Product getProduct(@PathParam("id") String id) {
    // ... logic that might be slow
}
```

- Throws a `TimeoutException` on timeout.

@Retry

- **Problem:** An operation fails due to a transient issue (e.g., temporary network glitch).
- **Solution:** Automatically retry the operation a few times.

```
@GET
@Path("/{id}")
@Retry(maxRetries = 3, delay = 200) // Retry up to 3 times, with 200ms delay
public Station getStationById(@PathParam("id") String id) {
    // ... call to a potentially flaky service
}
```

- Highly configurable (`maxDuration` , `jitter` , `retryOn` , `abortOn`).

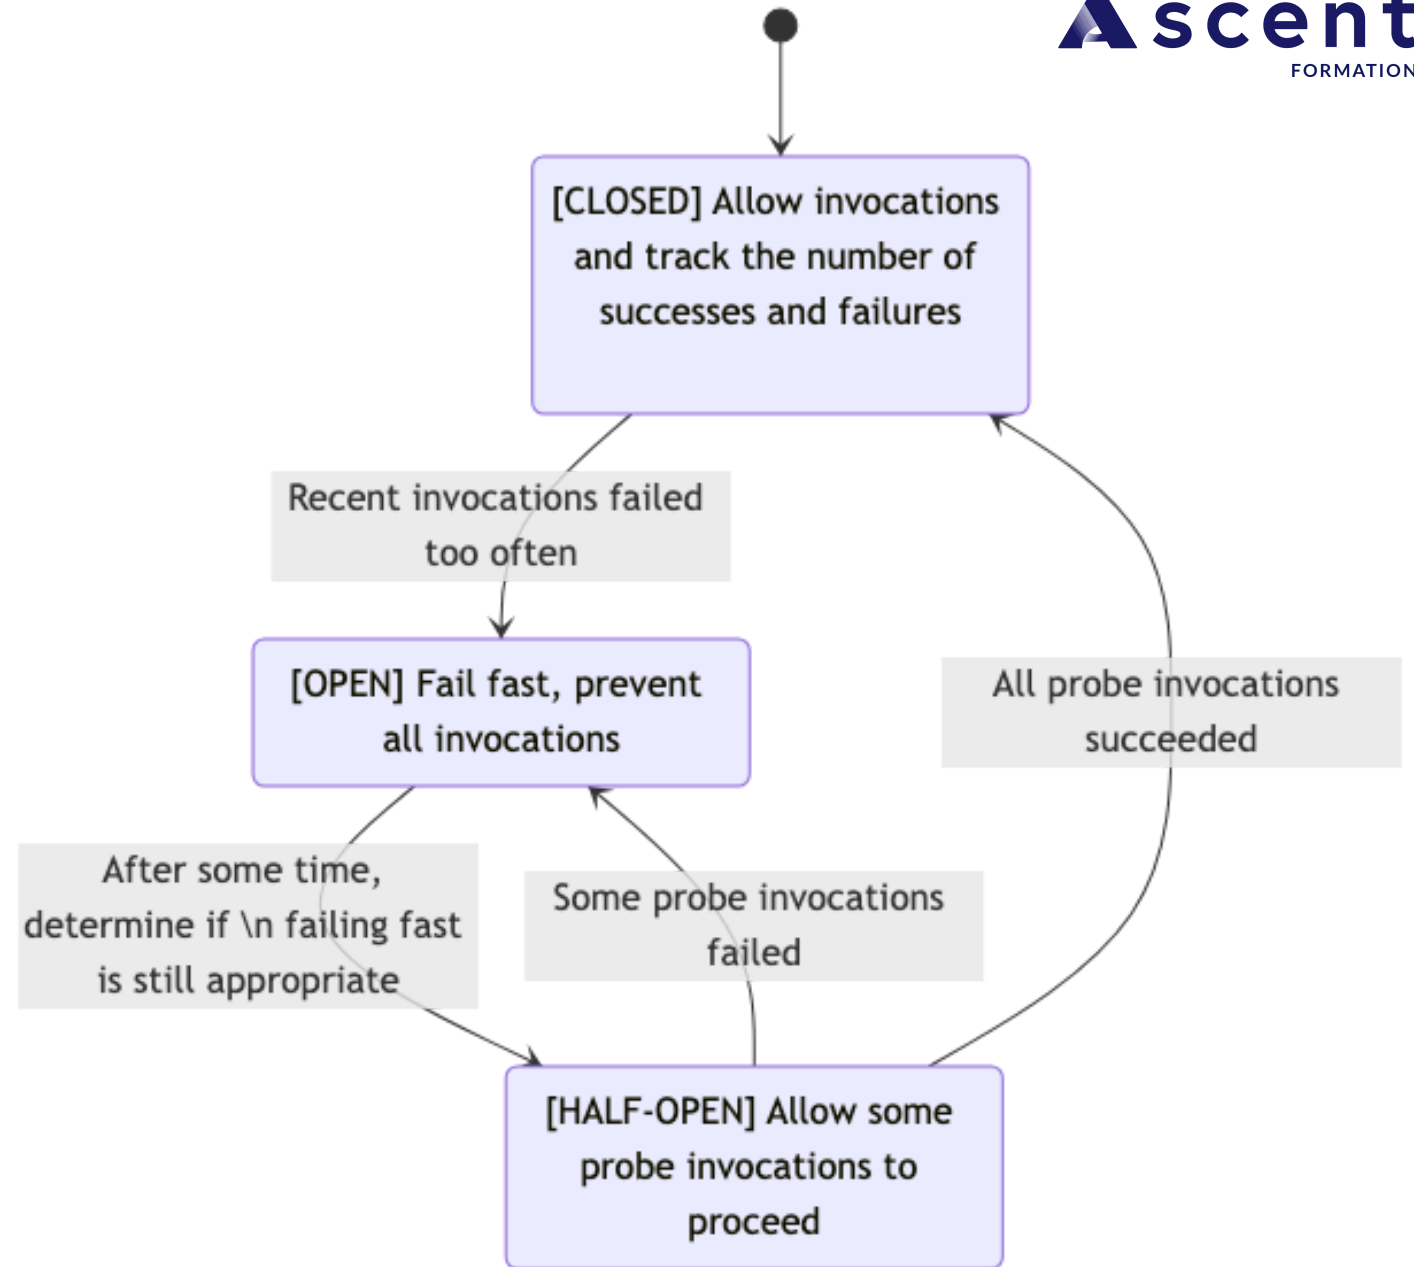
@CircuitBreaker

- **Problem:** A downstream service is down. Repeatedly retrying is wasteful and can make recovery harder (thundering herd).
- **Solution:** "Open" a circuit to fail-fast, and periodically check if the service has recovered.

States:

1. **CLOSED:** Operations are executed. If failures exceed a threshold, moves to OPEN.
2. **OPEN:** Operations fail immediately (`CircuitBreakerOpenException`). After a delay, moves to HALF-OPEN.
3. **HALF-OPEN:** A trial request is allowed. If it succeeds, moves to CLOSED. If it fails, moves back to OPEN.

@CircuitBreaker State Machine



@CircuitBreaker Example

```
@GET
@Path("/{id}")
@CircuitBreaker(
    requestVolumeThreshold = 10, // in the last 10 requests...
    failureRatio = 0.6,          // ...if 60% fail...
    delay = 5000,                // ...open the circuit for 5s
    successThreshold = 2         // need 2 successes to close
)
public Station getStationById(@PathParam("id") String id) {
    // ... call to a service that might be down
}
```

[CircuitBreaker SmallRye Documentation](#)

@Fallback

- **Problem:** An operation has failed (after timeouts and retries), but we want to provide a graceful response instead of an exception.
- **Solution:** Specify a method to call to get an alternative result.

```
@GET
@Path("/{id}")
@Retry(maxRetries = 2)
@Fallback(fallbackMethod = "getUnknownStation")
public Station getStationById(@PathParam("id") String id) {
    // ... call to a failing service
}
```

@Fallback

- The fallback method must have a compatible signature.

```
public Station getUnknownStation(String id) {  
    // Return a default or cached value  
    return new Station(id, "Unknown Station");  
}
```

@Fallback Configuration

You can control precisely when the fallback is triggered.

- **applyOn** : Specifies an array of exceptions that should trigger the fallback. If empty, all exceptions trigger it.
- **skipOn** : Specifies an array of exceptions that should *not* trigger the fallback.

@Fallback Configuration


```
@POST
@Fallback(
    fallbackMethod = "handleCreationFailure",
    applyOn = {WebApplicationException.class}, // Only apply on web exceptions
    skipOn = {IllegalArgumentException.class} // But not for bad arguments
)
public Response create(TrainStop stop) {
    // ... logic that might fail
}
```

This gives you fine-grained control over your error handling strategy.

[SmallRye Fallback Documentation](#)

@Bulkhead

- **Problem:** A single slow dependency can exhaust all available resources, causing the entire application to become unresponsive.
- **Analogy:** The bulkhead pattern is named after the partitioned sections of a ship's hull. If one compartment floods, the bulkheads prevent the entire ship from sinking.
- **Solution:** Isolate failures in one part of the system from affecting others by limiting the number of concurrent requests to a specific service.



@Bulkhead in Quarkus

- **Semaphore-based:** Limits the number of concurrent requests to a method.
- If the limit is reached, subsequent calls will fail immediately with a `BulkheadException`.

```
@GET
@Path("/{id}")
@Bulkhead(5) // Only 5 concurrent requests allowed
public Station getStationById(@PathParam("id") String id) {
    // ... call to a service that can be slow
}
```

- **Thread-pool-based (for async):** Provides a dedicated thread pool for the annotated method, isolating its execution from other threads.

[SmallRye Bulkhead Documentation](#)

@RateLimit

- **Problem:** A client is calling an endpoint too frequently, potentially overloading the service or a downstream resource (like a message queue).
- **Solution:** Restrict the number of requests allowed within a specific time window.
- **Implementation:** This is a SmallRye-specific annotation, **not** part of the MicroProfile Fault Tolerance specification.

[RateLimit SmallRye Documentation](#)

@RateLimit in Quarkus

```
import io.smallrye.faulttolerance.api.RateLimit;
import java.time.temporal.ChronoUnit;

@POST
@Path("/notify")
@RateLimit(
    value = 10,
    window = 1,
    windowUnit = ChronoUnit.MINUTES
) // Allow only 10 requests per minute
public Response notify(String message) {
    // Publish message to a queue
    messageProducer.send(message);
    return Response.accepted().build();
}
```

- If the rate limit is exceeded, a `RateLimitException` is thrown.

Priority of Fault Tolerance Strategies

```
Fallback(  
  Retry(  
    CircuitBreaker(  
      RateLimit(  
        Timeout(  
          Bulkhead(  
            ... the guarded method ...  
          )  
        )  
      )  
    )  
  )  
)
```

Interactions with Other Strategies

Lunch Break (12:00 - 13:00)


Afternoon Session 1 (13:00 - 14:30)

Topic: Asynchronous Messaging


Why Asynchronous Messaging?

- **Decoupling:** Services don't need to know about each other directly. They just produce/consume messages.
- **Resilience:** The system can tolerate temporary unavailability of a consumer. Messages will be processed when it comes back online.
- **Scalability:** You can have multiple instances of a consumer processing messages from a queue in parallel.
- **Responsiveness:** Fire-and-forget. The client doesn't have to wait for the full processing to complete.

- Interaction between services via REST APIs

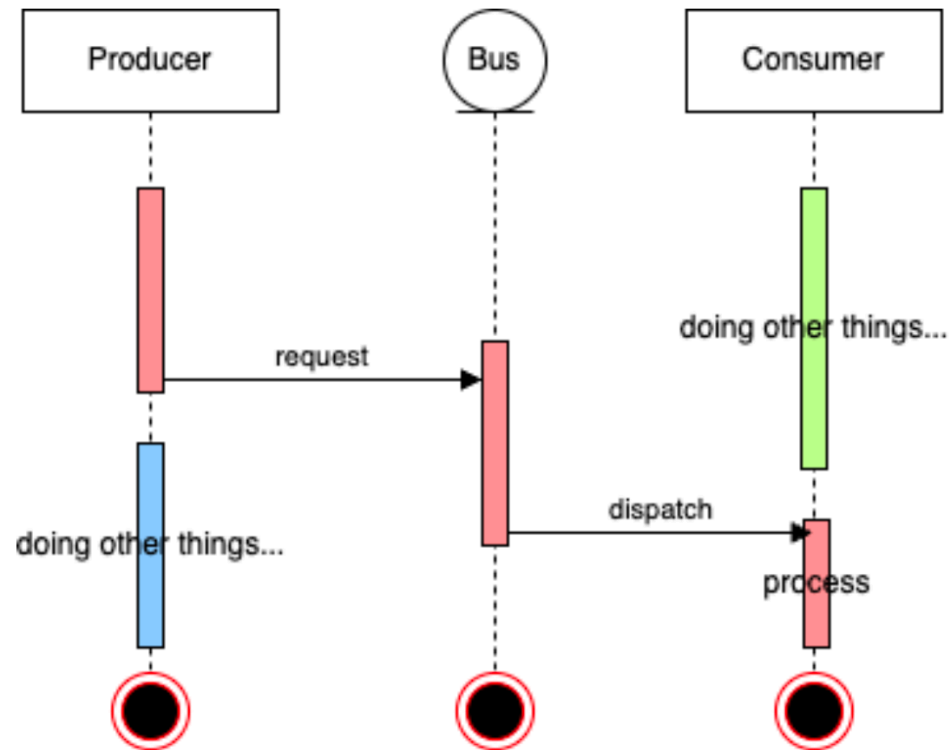


- Add a middleware to handle the communication
- Decouples the different services



- **Fire and forget**

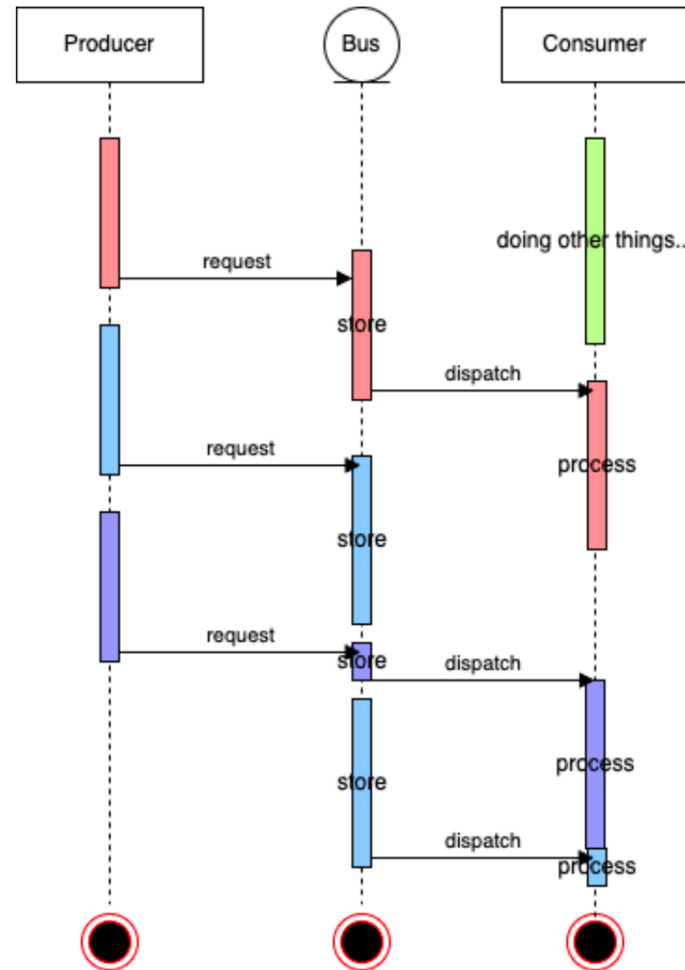
- In order for the producer to send a message, there is no need to wait for the receiver to be available



- Store and forward approach
- Message redelivery
- Message broker Redundancy



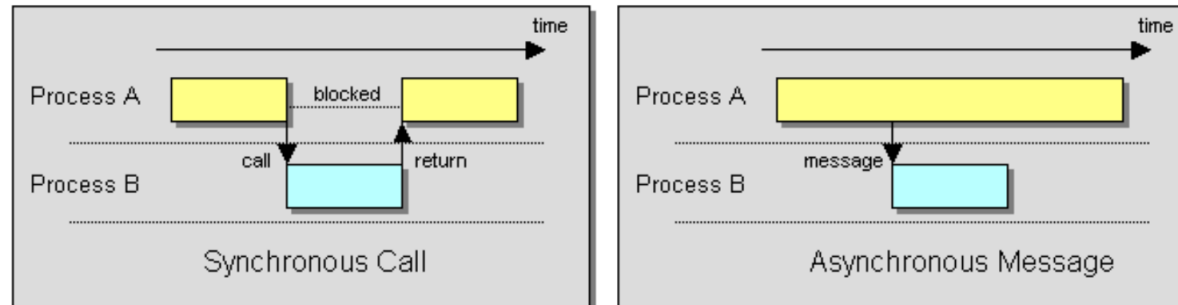
- Again, fire and forget
- Handle multiple requests at the same time



]

Synchronous vs. Asynchronous

- **Synchronous: Sender waits for a response to the message/request**
 - Request-response, gRPC
 - REST API call (GET, POST, etc)
 - SQL UPDATE
- **Asynchronous: Sender does not wait for receiver**
- **For request-response scenarios the sender provides a "callback"**
 - Javascript Promise/await/async (AJAX for oldschoolers)
 - Java/Guava Threads and Futures
 - C# await/async/Tasks



MicroProfile Reactive Messaging

A standard for building event-driven, asynchronous applications.

- **Core Concepts:**
 - **Channel:** A virtual, named destination. Can be in-memory or mapped to a message broker (like Kafka or AMQP).
 - **@Outgoing** : A method that produces messages and sends them to a channel.
 - **@Incoming** : A method that consumes messages from a channel.
 - **Connector:** The "glue" that maps a channel to a specific messaging technology.
- **Guide:** [Reactive Messaging](#)

Producer (@Outgoing)

A method that produces a stream of messages.

```
@ApplicationScoped
public class MessageGeneratorBean {

    @Outgoing("sink")
    public Multi<String> generate() {
        return Multi.createFrom().items("a", "b", "c");
    }

}
```

Consumer (@Incoming)

A method that is called for each message received from a channel.

```
@ApplicationScoped
public class MessageProcessingBean {

    @Incoming("source")
    public void process(String consumedPayload) {
        // process the payload
        consumedPayload.toUpperCase();
    }

}
```


Processor (@Incoming + @Outgoing)

A method that consumes from one channel, transforms the message, and sends it to another channel.

```
@Incoming("requests")
@Outgoing("quotes")
@Blocking
public Quote process(String quoteRequest) throws InterruptedException {
    // simulate some hard-working task
    Thread.sleep(200);
    return new Quote(quoteRequest, random.nextInt(100));
}
```

Connectors & Configuration

Connectors map channels to a message broker. Quarkus has connectors for Kafka, AMQP (RabbitMQ), MQTT, and more.

`application.properties`

```
amqp-host=amqp
amqp-port=5672
amqp-username=my-username
amqp-password=my-password

mp.messaging.incoming.prices.connector=smallrye-amqp
...
```

Afternoon Session 2 (14:45 - 16:00)

Activity: Lab on Asynchronous Messaging (Optional/Demo)

- <https://quarkus.io/guides/amqp>
- <https://quarkus.io/guides/amqp-reference#configuring-the-amqp-broker-access>
- <https://quarkus.io/guides/messaging-virtual-threads>

End of Day 3

- Recap & Q&A
- Preview of Day 4:
 - i. Observability
 - ii. Metrics
 - iii. Azure Service Bus Messaging
 - iv. Reactive Programming vs Virtual Threads
 - v. Transactions