

## **Morning Session 1 (09:00 - 11:00)**

**Topic: Lab 7 - The Resilient SAGA (Hands-on)**

# The SAGA Pattern (Theory)

- **What is a SAGA?** A sequence of local transactions where each transaction updates data within a single service.
- **Purpose:** To manage data consistency across microservices without using distributed transactions.
- **Long-Lived Transactions:** Ideal for processes that span multiple services and may take a long time to complete.
- **Failure Handling:** If a local transaction fails, the SAGA executes compensating transactions to undo the preceding transactions.

# Lab 7, Part 1: The SAGA Trigger

## The API Endpoint

- **Objective:** Create the initial API endpoint that starts the SAGA.
- **Action:** The endpoint receives a request and produces the first event/message.
- **Benefit:** Decouples the synchronous API call from the asynchronous business process.

# Lab 7, Part 2: The SAGA Step & Resilience

## The Consumer

- **Objective:** Implement a message consumer for a step in the SAGA.
- **Action:** The consumer processes the message, executes a local transaction, and applies resilience patterns.
- **Resilience:** Use `@Retry` to handle transient failures during message processing.

# Lab 7, Part 3: The Choreography

## Completing the SAGA

- **Objective:** Connect the SAGA steps by having consumers produce subsequent events.
- **Action:** Upon successful processing, the consumer sends a new message to trigger the next step in the workflow.
- **Choreography:** Services communicate directly with each other via events without a central controller.

**Morning Break (11:00 - 11:15)**

## **Morning Session 2 (11:15 - 12:00)**

**Topic: SAGA and Messaging Discussion**

# SAGA: Choreography vs. Orchestration

- **Choreography (Our Approach):**
  - No central coordinator.
  - Each service produces/listens to events and decides what to do.
  - **Pros:** Simple, loosely coupled.
  - **Cons:** Hard to track the process, risk of cyclic dependencies.
- **Orchestration:**
  - A central "SAGA orchestrator" tells services what to do.
  - The orchestrator manages the state of the entire transaction.
  - **Pros:** Centralized logic, easier to monitor.
  - **Cons:** Tighter coupling, orchestrator can become a bottleneck.

# Compensating Transactions & DLQs

- **Compensating Transactions:**
  - The "undo" operation for a SAGA step.
  - If step  $T_{-i}$  fails, a compensating transaction  $C_{-i-1}$  is run to reverse the effects of  $T_{-i-1}$ .
  - Essential for maintaining data consistency when things go wrong.
- **Dead-Letter Queues (DLQ):**
  - A dedicated queue for messages that could not be processed successfully (after retries).
  - Prevents "poison pills" from blocking the main queue.
  - Allows for manual inspection and reprocessing of failed messages.

# Key Messaging Concepts

- **Idempotency:** Can a message be processed multiple times with the same result?  
Consumers *must* be idempotent.
- **"At-Least-Once" Delivery:** Most message brokers guarantee this. Your system must handle duplicates.
- **Message Ordering:** Is the order of messages guaranteed? Often, it's only guaranteed within a partition.
- **Poison Pills:** A malformed or problematic message that repeatedly causes a consumer to fail.

# SmallRye Reactive Messaging

- **Quarkus's Solution:** Implements the MicroProfile Reactive Messaging specification.
- **@Incoming / @Outgoing :** Annotations to define consumers and producers.
- **Channels:** Logical names for message streams, configured in `application.properties` .
- **Connectors:** The "glue" that connects a channel to a specific broker (Kafka, AMQP, etc.).

**Lunch Break (12:00 - 13:00)**

## **Afternoon Session 1 (13:00 - 14:00)**

**Topic: Observing Your Microservices**

# Why Observability?

In a complex microservices architecture, you need to answer:

- **Is the service running?** (Health Checks)
- **How is the service performing?** (Metrics)
- **What happened during a specific request?** (Distributed Tracing - covered later)

Observability is key to debugging, monitoring, and maintaining a healthy system.

# MicroProfile Health

Exposes health check procedures to let orchestrators (like Kubernetes) know if your application is healthy.

- Two types of checks:
  - Liveness (`/q/health/live`): Is the application running? If this fails, the container should be restarted.
  - Readiness (`/q/health/ready`): Is the application ready to accept requests? If this fails, the container should be temporarily removed from the load balancer.
- Guide: [MicroProfile Health](#)

# Implementing Health Checks

Create a bean that implements `HealthCheck` and is annotated with `@Liveness` or `@Readiness`.

```
@Liveness
@ApplicationScoped
public class DatabaseConnectionHealthCheck implements HealthCheck {

    @Override
    public HealthCheckResponse call() {
        if (isDatabaseConnectionOk()) {
            return HealthCheckResponse.up("Database connection is OK");
        } else {
            return HealthCheckResponse.down("Database connection is down");
        }
    }
}
```

# Lab 8: Observability (Readiness)

- **Objective:** Implement readiness checks so Kubernetes knows if the service is ready to consume messages.
- **Verify Connections:**
  - Connection to the PostgreSQL Database.
  - Connection to the Azure Service Bus.
- **Demo:** Show the service reporting "DOWN" when the broker is inaccessible and "UP" when it's available.

**Afternoon Break (14:00 - 14:15)**

## **Afternoon Session 2 (14:15 - 15:45)**

**Topic: Observability & Production Operations**

# 1. Metrics with Micrometer

(Prometheus & Grafana)

- **Topic:** Instrumenting the application for performance monitoring.
- **Micrometer:** The de facto standard for metrics in the Java ecosystem.
- **Demo Steps:**
  - i. Add `quarkus-micrometer-registry-prometheus`.
  - ii. Add a `@Counted` metric to the SAGA consumer.
  - iii. Show the new custom metric in the `/q/metrics` output.

# Custom Metrics with `@Counted`

Easily count how many times a method is invoked. This is a great way to monitor when a downstream service is failing.

```
@ApplicationScoped
public class StationServiceFallbackHandler {

    @Counted(name = "station_service_fallbacks_total", description = "Counts total fallbacks for the Station Service.")
    public Station getStationByIdFallback(String id) {
        Log.warnf("Fallback for getStationById, station %s not found.", id);
        return new Station("0", "Station Details Currently Unavailable");
    }
}
```

- This metric will be exposed in Prometheus format.

# Conceptual Grafana Dashboard

Visualize your metrics to gain insights into application performance.

!Grafana Dashboard, bg right:70% height:80%

## 2. Centralized & Structured Logging

- **Topic:** Preparing logs for production analysis (e.g., Splunk, Loki).
- **Structured Logging:** Writing logs in a machine-readable format like JSON.
- **Demo Steps:**
  - i. Enable JSON logging in `application.properties`.
  - ii. Add contextual data (`MDC.put(...)`) to the SAGA consumer.
  - iii. Show the resulting JSON log enriched with the ID.

## 3. Distributed Tracing (OpenTelemetry)

- **Topic:** Tracing a single request across all services and message brokers.
- **OpenTelemetry:** A standard for generating and collecting telemetry data (traces, metrics, logs).
- **The "Aha!" Moment:**
  - Add `quarkus-opentelemetry`.
  - Trigger the SAGA and show that the same `traceId` is automatically propagated from the initial API call, across Azure Service Bus, and appears in the consumer's log.
- **Discussion:** How this enables end-to-end tracing in tools like Jaeger or Datadog.

## **Afternoon Session 3 (15:45 - 16:15)**

**Topic: Reactive vs. Virtual Threads**

# 1. Reactive Programming (Mutiny)

- **What is Reactive?** An asynchronous programming paradigm concerned with data streams and the propagation of change.
- **Key Concepts:**
  - **Non-blocking:** Threads are not held waiting for I/O operations.
  - **Event-driven:** The program flow is determined by events.
  - **Back-pressure:** A mechanism for consumers to control the rate of production.
- **Connection:** This is the engine behind SmallRye Reactive Messaging  
(`@Incoming` / `@Outgoing`).
- **Demo:** Show a simple Mutiny `Uni` or `Multi` to demonstrate the "pipeline" concept.

## 2. Virtual Threads (Project Loom)

- **What are Virtual Threads?** Lightweight threads managed by the JVM, not the OS.
- **Key Concepts:**
  - Drastically increases the number of concurrent operations a server can handle.
  - Allows developers to write simple, sequential, blocking-style code that scales massively.
- **Connection:** The "new" way to achieve high concurrency without the complexity of reactive programming.
- **Demo:** Show a JAX-RS endpoint annotated with `@RunOnVirtualThread` and discuss the performance implications.

# Open Discussion & Feedback (16:15 - 16:30)

- Course Feedback & Questionnaire
- Open Discussion
- Resources

# End of Training

Thank You!



**Scott Messner**

-  Email [scott.m.messner@gmail.com](mailto:scott.m.messner@gmail.com)
-  GitHub [github.com/blackstrype/quarkus-microservices-training](https://github.com/blackstrype/quarkus-microservices-training)
-  LinkedIn [linkedin.com/in/scott-messner-0264231a](https://linkedin.com/in/scott-messner-0264231a)