

Quarkus Microservices

Day 4

- Reactive Programming / Virtual Threads
- SAGA and Messaging
- Observability
- Production Operations

Morning Session 1 (09:00 - 09:45)

Topic: Reactive vs. Virtual Threads

1. Reactive Programming (Mutiny)

- **What is Reactive?** An asynchronous programming paradigm concerned with data streams and the propagation of change.
- **Key Concepts:**
 - **Non-blocking:** Threads are not held waiting for I/O operations.
 - **Event-driven:** The program flow is determined by events.
 - **Back-pressure:** A mechanism for consumers to control the rate of production.
- **Connection:** This is the engine behind SmallRye Reactive Messaging
(`@Incoming` / `@Outgoing`).
- **Demo:** Show a simple Mutiny `Uni` or `Multi` to demonstrate the "pipeline" concept.
 - <https://quarkus.io/guides/getting-started-reactive>

1. Reactive Programming Fork-Join

```
public Uni<DashboardData> getDashboard(String userId) {  
    // 1. Define the two independent, asynchronous actions  
    Uni<User> userUni = userService.getUser(userId);  
    Uni<List<Notification>> notificationsUni = notificationService.getNotifications(userId);  
  
    // 2. Combine them  
    return Uni.combine().all().unis(userUni, notificationsUni)  
        // 3. The 'combine' operation returns a List<Object>  
        .asTuple()  
  
        // 4. When both are complete, transform the tuple into the final DTO  
        .onItem().transform(tuple -> {  
            User user = tuple.getItem1();  
            List<Notification> notifications = tuple.getItem2();  
            return new DashboardData(user, notifications);  
        });  
}
```

1. Reactive vs. Completable Futures

```
public class DashboardService {  
  
    private final UserService userService = new UserService();  
    private final NotificationService notificationService = new NotificationService();  
  
    // You need a thread pool to run tasks in parallel.  
    // Use a try-with-resources block or manage its lifecycle carefully.  
    public DashboardData getDashboard(String userId)  
        throws ExecutionException, InterruptedException {  
  
        // Use a fixed thread pool for parallel execution  
        try (ExecutorService executor = Executors.newFixedThreadPool(2)) {  
  
            // 1. Fork: Submit the first task to the thread pool  
            CompletableFuture<User> userFuture = CompletableFuture.supplyAsync(() ->  
                userService.getUser(userId), executor  
            );  
  
            // 2. Fork: Submit the second task to the thread pool (runs in parallel)  
            CompletableFuture<List<Notification>> notificationsFuture = CompletableFuture.supplyAsync(() ->  
                notificationService.getNotifications(userId), executor  
            );  
  
            // 3. Join & Combine:  
            //     thenCombine waits for *both* futures to complete,  
            //     then provides their results to the combining function.  
            CompletableFuture<DashboardData> combinedFuture = userFuture.thenCombine(  
                notificationsFuture,  
                (user, notifications) -> new DashboardData(user, notifications)  
            );  
  
            // 4. Get Result (Blocking):  
            //     This is the main difference from Mutiny.  
            //     We must now block the current thread to wait for the result.  
            return combinedFuture.get();  
        }  
    }  
}
```

2. Virtual Threads (Project Loom)

- **What are Virtual Threads?** Lightweight threads managed by the JVM, not the OS.
- **Key Concepts:**
 - Drastically increases the number of concurrent operations a server can handle.
 - Allows developers to write simple, sequential, blocking-style code that scales massively.
- **Connection:** The "new" way to achieve high concurrency without the complexity of reactive programming.
- **Demo:** Quotes Processor using `@RunOnVirtualThread`

Morning Session 2 (09:45 - 10:45)

Topic: SAGA and Messaging In Practice

The SAGA Pattern (Introduction)

- **What is a SAGA?** A sequence of local transactions where each transaction updates data within a single service.
- **Purpose:** To manage data consistency across microservices without using distributed transactions.
- **Long-Lived Processes:** Ideal for workflows that span multiple services and may take a long time to complete.
- **Failure Handling:** If a local transaction fails, the SAGA executes compensating transactions to undo the preceding transactions.

Lab 7, Part 1: The SAGA Trigger

The API Endpoint

- **Objective:** Create the initial API endpoint that starts the SAGA.
- **Action:** The endpoint receives a request and produces the first event/message.
- **Benefit:** Decouples the synchronous API call from the asynchronous business process.

Morning Break (11:00 - 11:15)

Morning Session 2.2 (11:15 - 12:00)

Topic: Lab 7 - The Resilient SAGA (Hands-on)

Lab 7, Part 2: The SAGA Step & Resilience

The Consumer

- **Objective:** Implement a message consumer for a step in the SAGA.
- **Action:** The consumer processes the message, executes a local transaction, and applies resilience patterns.
- **Resilience:** Use `@Retry` to handle transient failures during message processing.

Lab 7, Part 3: The Choreography

Completing the SAGA

- **Objective:** Connect the SAGA steps by having consumers produce subsequent events.
- **Action:** Upon successful processing, the consumer sends a new message to trigger the next step in the workflow.
- **Choreography:** Services communicate directly with each other via events without a central controller.

Lunch Break (12:00 - 13:00)

Afternoon Session 1 (13:00 - 14:00)

Topic: SAGA and Messaging Discussion

SAGA: Choreography vs. Orchestration

- **Choreography (Our Approach):**
 - No central coordinator.
 - Each service produces/listens to events and decides what to do.
 - **Pros:** Simple, loosely coupled.
 - **Cons:** Hard to track the process, risk of cyclic dependencies.
- **Orchestration:**
 - A central "SAGA orchestrator" tells services what to do.
 - The orchestrator manages the state of the entire transaction.
 - **Pros:** Centralized logic, easier to monitor.
 - **Cons:** Tighter coupling, orchestrator can become a bottleneck.

SAGA in Quarkus: Narayana LRA

- MicroProfile LRA Spec
- Narayana LRA Participant Support



SAGA: Choreography

Choreography SAGA Diagram

SAGA: Orchestration

 Orchestration SAGA Diagram

Compensating Transactions & DLQs

- **Compensating Transactions:**

- The "undo" operation for a SAGA step.
- If step T_{-i} fails, a compensating transaction c_{-i-1} is run to reverse the effects of T_{-i-1} .
- Essential for maintaining data consistency when things go wrong.

- **Dead-Letter Queues (DLQ):**

- A dedicated queue for messages that could not be processed successfully (after retries).
- Prevents "poison pills" from blocking the main queue.
- Allows for manual inspection and reprocessing of failed messages.

Key Messaging Concepts

- **Idempotency:** Can a message be processed multiple times with the same result?
Consumers *must* be idempotent.
- **"At-Least-Once" Delivery:** Most message brokers guarantee this. Your system must be designed to handle duplicates.
- **Message Ordering:** Is the order of messages guaranteed? Often, it's only guaranteed within a partition.
- **Poison Pills:** A malformed or problematic message that repeatedly causes a consumer to fail.

SmallRye Reactive Messaging

- **Quarkus's Solution:** Implements the MicroProfile Reactive Messaging specification.
- **@Incoming / @Outgoing** : Annotations to define consumers and producers.
- **Channels:** Logical names for message streams, configured in `application.properties` .
- **Connectors:** The "glue" that connects a channel to a specific broker (Kafka, AMQP, etc.).

Afternoon Session 2 (14:00 - 15:00)

Topic: Observing Your Microservices

Why Observability?

In a complex microservices architecture, you need to answer:

- **Is the service running?** (Health Checks)
- **How is the service performing?** (Metrics)
- **What happened during a specific request?** (Distributed Tracing - covered later)

Observability is key to debugging, monitoring, and maintaining a healthy system.

MicroProfile Health

Exposes health check procedures to let orchestrators (like Kubernetes) know if your application is healthy.

- **Three types of checks:**
 - **Liveness (`/q/health/live`)**: Is the application running? If this fails, the container should be restarted.
 - **Readiness (`/q/health/ready`)**: Is the application ready to accept requests? If this fails, the container should be temporarily removed from the load balancer.
 - **Startup Check (`/q/health/started`)**: Is the app component started ? Allows for a DOWN status while starting up.
- **Guide:** [MicroProfile Health](#)

Implementing Health Checks

Create a bean that implements `HealthCheck` and is annotated with `@Liveness` or `@Readiness`.

```
@Liveness
@ApplicationScoped
public class DatabaseConnectionHealthCheck implements HealthCheck {

    @Override
    public HealthCheckResponse call() {
        if (isDatabaseConnectionOk()) {
            return HealthCheckResponse.up("Database connection is OK");
        } else {
            return HealthCheckResponse.down("Database connection is down");
        }
    }
}
```

Lab 8: Observability (Readiness)

- **Objective:** Implement readiness checks so Kubernetes knows if the service is ready to consume messages.
- **Verify Connections:**
 - Connection to the PostgreSQL Database.
 - Connection to the Message Broker.
- **Demo:** Show the service reporting "DOWN" when the broker is inaccessible and "UP" when it's available.

Afternoon Break (15:00 - 15:15)

Afternoon Session 3 (15:15 - 16:15)

Topic: Observability & Production Operations

1. Metrics with Micrometer

(Prometheus & Grafana)

- **Topic:** Instrumenting the application for performance monitoring.
- **Micrometer:** The de facto standard for metrics in the Java ecosystem.
- **Steps:**
 - i. Add `quarkus-smallrye-metrics`.
 - ii. Add a `@Counted` metric to the SAGA consumer.
 - iii. Show the new custom metric in the `/q/metrics` output.
 - iv. Add `quarkus-micrometer-registry-prometheus` to export in a Prometheus format

Custom Metrics with `@Counted`

Easily count how many times a method is invoked. This is a great way to track how often a downstream service is failing.

```
@ApplicationScoped
public class StationServiceFallbackHandler {

    @Counted(name = "station_service_fallbacks_total", description = "Counts total fallbacks for the Station Service.")
    public Station getStationByIdFallback(String id) {
        Log.warnf("Fallback for getStationById, station %s not found.", id);
        return new Station("0", "Station Details Currently Unavailable");
    }
}
```

- This metric will be exposed in Prometheus format.

Conceptual Grafana Dashboard

Visualize your metrics to gain insights into application performance.



2. Centralized & Structured Logging

- **Topic:** Preparing logs for production analysis (e.g., Splunk, Loki).
- **Structured Logging:** Writing logs in a machine-readable format like JSON.
- **Demo Steps:**
 - i. Enable JSON logging in `application.properties`.
 - ii. Add contextual data (`MDC.put(...)`) to the SAGA consumer.
 - iii. Show the resulting JSON log enriched with the contextual ID.

3. Distributed Tracing (OpenTelemetry)

- **Topic:** Tracing a single request across all services and message brokers.
- **OpenTelemetry:** A standard for generating and collecting telemetry data (traces, metrics, logs).
- **The "Aha!" Moment:**
 - Add `quarkus-opentelemetry`.
 - Trigger the SAGA and show that the **same** `traceId` is automatically propagated from the initial API call, across Azure Service Bus, and appears in the consumer's log.
- **Discussion:** How this enables end-to-end tracing in tools like Jaeger or Datadog.

Open Discussion & Feedback (16:15 - 16:30)

- Course Feedback & Questionnaire
- Open Discussion
- Resources

End of Training

Thank You!

-  Scott Messner l:50px **Scott Messner**
-  Email l:50px scott.m.messner@gmail.com
-  GitHub l:50px github.com/blackstrype/quarkus-microservices-training
-  LinkedIn l:50px linkedin.com/in/scott-messner-0264231a